


Machine Learning for Quality Assurance of Lutz Catchment Runoff Data

Gillian A. McGinnis, BA

University of Arizona College of Information Science
INFO 698 Capstone Project, Fall 2025



Background

Water runoff data has been continuously collected by the Smithsonian Tropical Research Institute (STRI) at the **Lutz Catchment weir** on Barro Colorado Island since 1972, & with electronic sensors **since 1989**. [1, 2]



Different "failure modes" impact data quality, resulting in the need for adjustments which are currently conducted manually.

Goals

Determine if models can be constructed to **effectively identify windows of failure** and eventually conduct **quality assurance** corrections on the raw runoff values without the need for manual intervention.

Challenges

➤ **New project & approach**
Past stochastic approaches had failed; extensive background research was necessary to determine appropriate machine learning model types

➤ **Gaps & missing data**
Sensor failures, missing values, and inconsistent labels, comments, & flags

➤ **Differences in data frequency**
Runoff & rain: *every 5 minutes for 36 years from 1 site*

Soil moisture: *once a week from 10 sites × 2 depths each*

➤ **Computational power**

More than 4.15 million total entries

➤ **Access & communication**


Primary data contact was unavailable the majority of Oct & Nov due to the federal government shutdown

Failure Modes

Five major failure modes cause data quality issues.

Blockage ★

Debris blocks the weir's 'V'
Fix: data pivot or decay curve
*Most difficult failure mode to identify & correct




Spike

Short & abrupt changes in level
Fix: smoothing to neighbors' values
*Differs from sudden increase in flows following heavy rain




Calibration

Misalignment with standard (x)
Fix: baseline correction
*Points used for corrections, but Blockage can render ineffective




Sub-Zero

Stream runs dry or is drained
Fix: setting to zero
*Data after rain may be unrecoverable



Signal Noise

Equipment failure
Fix: N/A
*Impossible to manually correct, and resists standard de-noising



Methods for Blockage flagging


Language: Python (Jupyter Notebooks)
Dataset: n = 3,472,682 annotated points

Process:

1. Introduce raw data (*imported from CSV*)
2. Prepare data (*clean & wrangle*)
3. Engineer features (*time, lags, & rolling stats*)
4. Remove high-correlation features (>0.97)
5. Conduct Train/Test split (80:20)
6. Tune XGBoost hyperparameters (*w/ PR-AUC—good for imbalanced classes*)
7. Fit to get out-of-fold predictions (OOF)
8. Tune post-hoc smoothing parameters w/ F1 (*median-windowing & classification threshold*)
9. Fit tuned XGBoost model to entire training set
10. Predict on the test set (*the held-out 20%*)
11. Analyze results (*w/ & w/o post-hoc params.*)

Splitting


Time series data should *not* be randomly split like typical k-fold cross-validation:
Chronologically-ordered full set



Splits to tune hyperparameters use expanding windows of time:



Once tuned, post-hoc parameters can be tuned using OOF predictions:



Once the model is fully tuned:



Training Results

n = 2,778,145
19 Jul 1989 11:55 -thru- 08 Mar 2018 21:50*

Number of features:

22 original Runoff, rain, and ×20 soil
+ 117 engineered Lags and rolling stats
- 31 high-correlated feats >0.97 correlated w/ another
= 108 input features

Hyperparameters:

n estimators 122 Number of trees
Learning rate 0.102 Impact of new trees
Max depth 3 Max depth of individual trees
Subsample 0.6455 Random subset of training rows when building each tree
Column " by tree 0.709 Random subset of features
Scale pos. weight 11 Handles class imbalance
Gamma 0.128 Minimum loss reduction to split
Alpha 0.936 L1 regularization


Post-hoc smoothing & tuning:

Median window size = 29 & threshold = 0.307
After removing marginal gains in scoring:
No median windowing & threshold = 0.317

Test Set Performance

n = 694,537
08 Mar 2018 21:55 -thru- 01 Aug 2025 13:00*

PR-curve of fitted model on the test set (shows performance across different thresholds)



	Accuracy	Precision	Recall	F1
Defaults	0.667	0.253	0.686	0.397
Win = 29	0.669	0.255	0.686	0.371
Win = 29 & Th = 0.307	0.509	0.204	0.845	0.329
Th = 0.317	0.518	0.206	0.838	0.331
Change	-0.149	-0.047	+0.152	-0.066


*Both sets are *not* perfectly continuous every 5 min due to gaps from occasional sensor failure, blips, & missing values

XGBoost: Extreme Gradient Boosting

XGBoost is a form of **gradient tree boosting**, which iteratively adds **weak learners** (*shallow, simple decision trees*) using gradient descent to minimize error.

It can also handle missing & incomplete data.

The final model is an **ensemble**, which sums up the predictions from all trees.



Interpretations

The model can identify most true blockages (*high recall*), but frequently has false alarms (*low precision*).

Threshold tuning improved model performance in identifying true blockages, but median window smoothing proved ineffective.

Limitations

The engineered features for XGBoost are **look-behind**, whereas manual corrections often use **look-ahead**. For example, a steep dropoff in runoff values followed by a calibration point (x) can indicate the end of a blockage.

Other failure modes (e.g., Spikes) may result in noisy data that the model struggles with interpreting.

Future Work

Create models for other failure modes. Not all modes will require models as complex as XGBoost due to reliance on fewer features, and it is likely that different classification algorithms may have to be considered for other failure mode detection models.

Address data correction automation. If a flagging model has poor performance, it may be necessary to make systems to apply adjustment models to manually-confirmed windows, since modifying accurate data harms quality.

Acknowledgements

Capstone advisor: Sriram Iyengar, PhD (*The University of Arizona College of Medicine, Phoenix*)

Data advisor, source, & access: Steven Paton, MSc (*Director of the Physical Monitoring Program, STRI*)

GPU assistance: Brian E. McGinnis, PhD

References

- [1] Barro Colorado (Clearing, Lutz, Conrad weir). Physical Monitoring.
- [2] Larsen, M. C.; Stallard, R. F.; Paton, S. Lutz Creek Watershed, Barro Colorado Island, Republic of Panama. *Hydrological Processes* 2021, 35 (4), e14157.
- [3] Introduction to Boosted Trees – xgboost 3.1.1 documentation.